



Introduction to Socket Programming




Protocol Families - TCP/IP

- Several protocols for different problems
- Protocol suites or protocol families: TCP/IP
- TCP/IP provides end-to-end connectivity specifying how data should be
 - Formatted,
 - Addressed,
 - Transmitted,
 - Routed, and
 - Received at the destination
- Can be used in the internet and in stand-alone private networks
- It is organized into layers

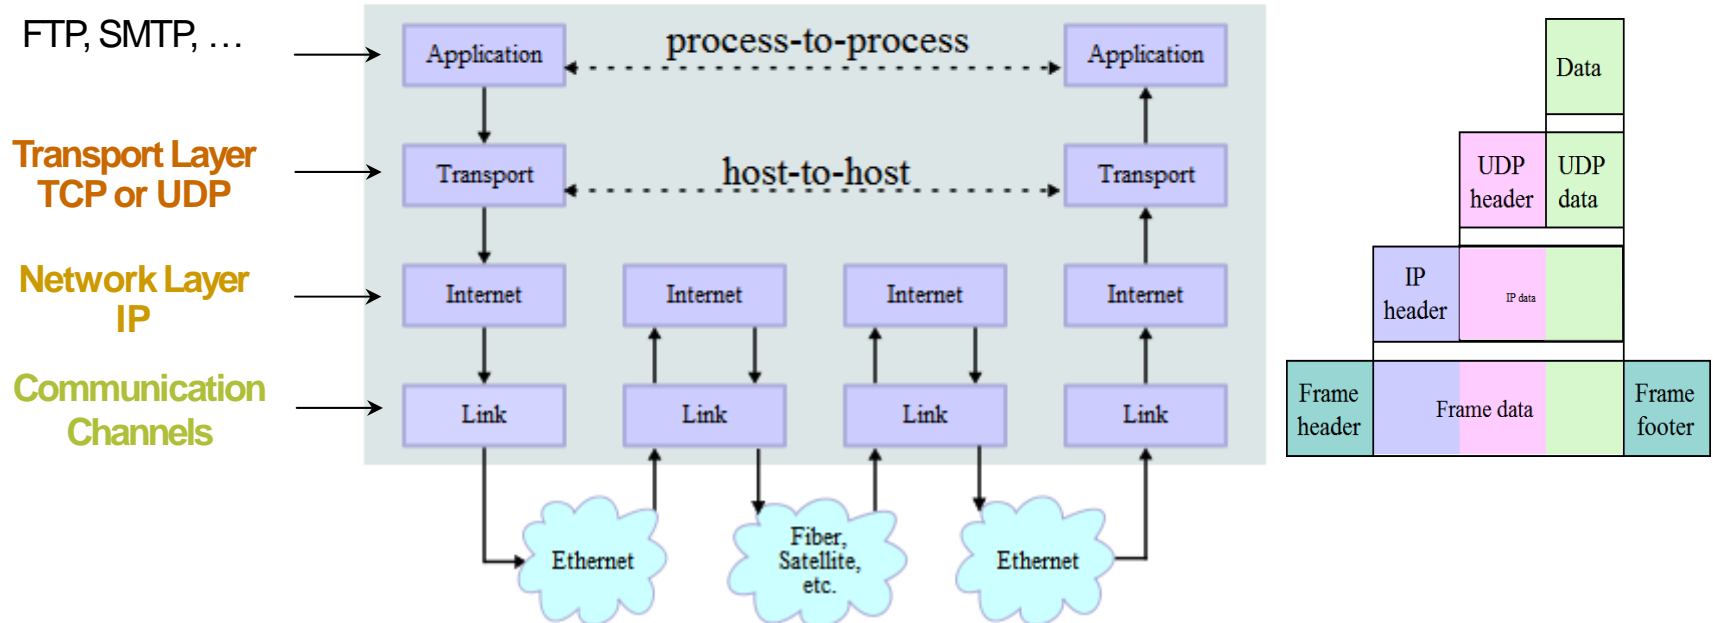
TCP/IP

Network Topology



*

Data Flow



* image is taken from "http://en.wikipedia.org/wiki/TCP/IP_model"

TCP vs UDP

- Both use **port numbers**
 - Application-specific construct serving as a communication endpoint
 - 16-bit unsigned integer, thus ranging from 0 to 65535 to provide end-to-end transport
- UDP: user datagram protocol
 - No acknowledgements and No retransmissions
 - Out of order, duplicates possible
 - Connectionless, i.e., App indicates destination for each packet
- TCP: transmission control protocol
 - Reliable **byte-stream channel** (in order, all arrive, no duplicates)
 - Similar to file I/O
 - Flow control
 - Connection-oriented
 - Bidirectional

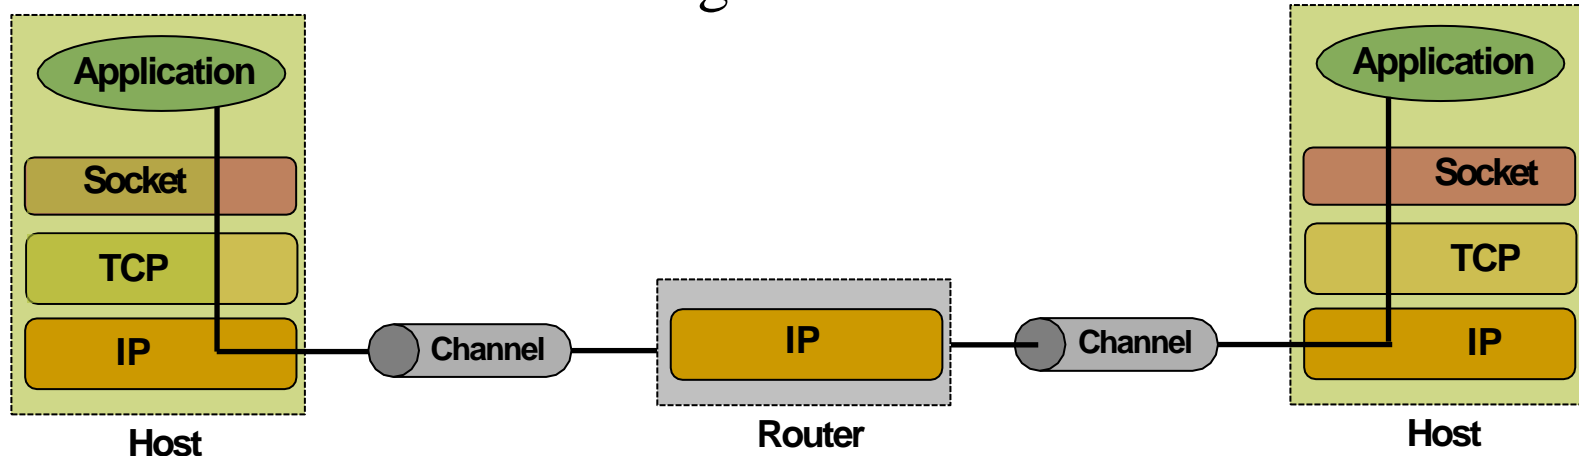
TCP vs UDP

- TCP is used for services with a large data capacity, and a persistent connection
- UDP is more commonly used for quick lookups, and single use query-reply actions.
- Some common examples of TCP and UDP with their default ports:

DNS lookup	UDP	53
FTP	TCP	21
HTTP	TCP	80
POP3	TCP	110
Telnet	TCP	23

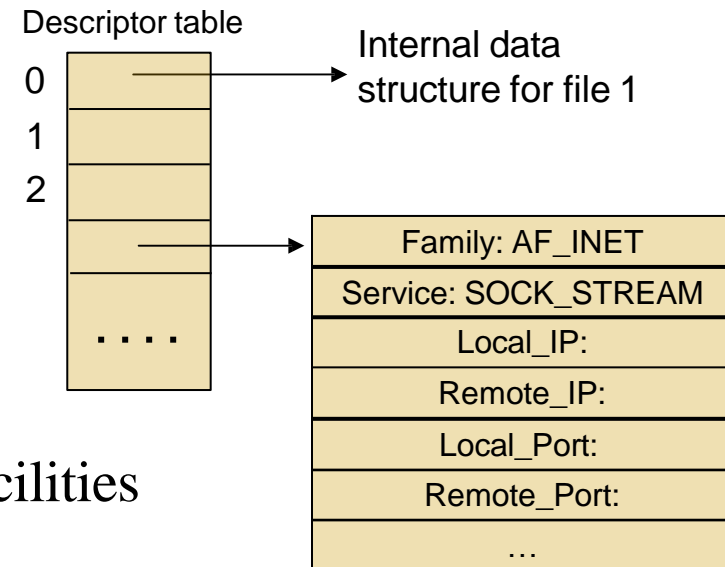
Berkley Sockets

- Universally known as Sockets
- It is an abstraction through which an application may send and receive data
- Provide generic access to interprocess communication services
 - e.g. IPX/SPX, Appletalk, TCP/IP
- Standard API for networking

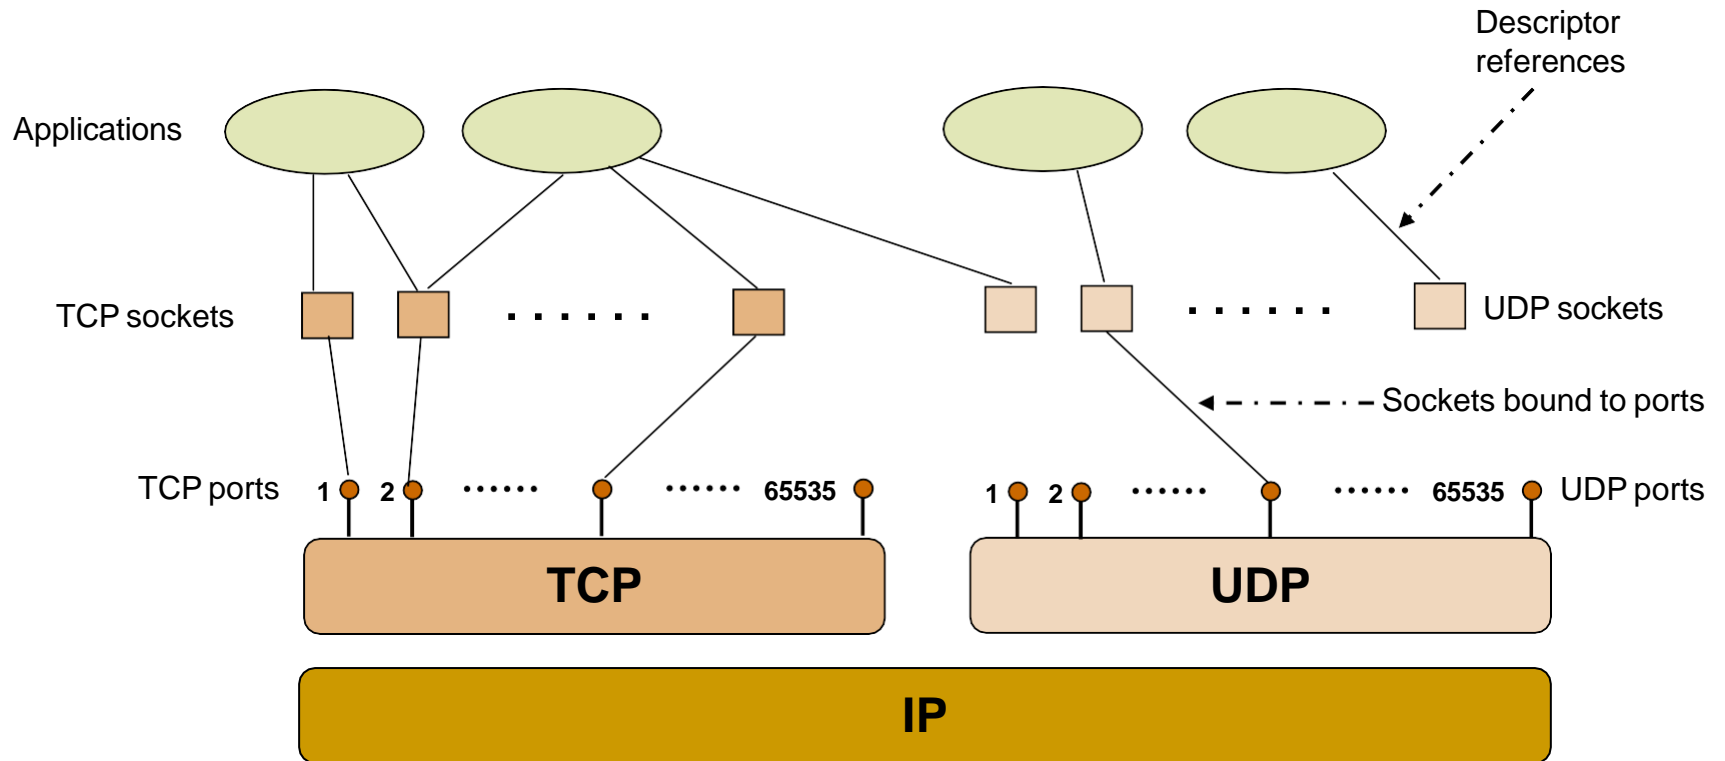


Sockets

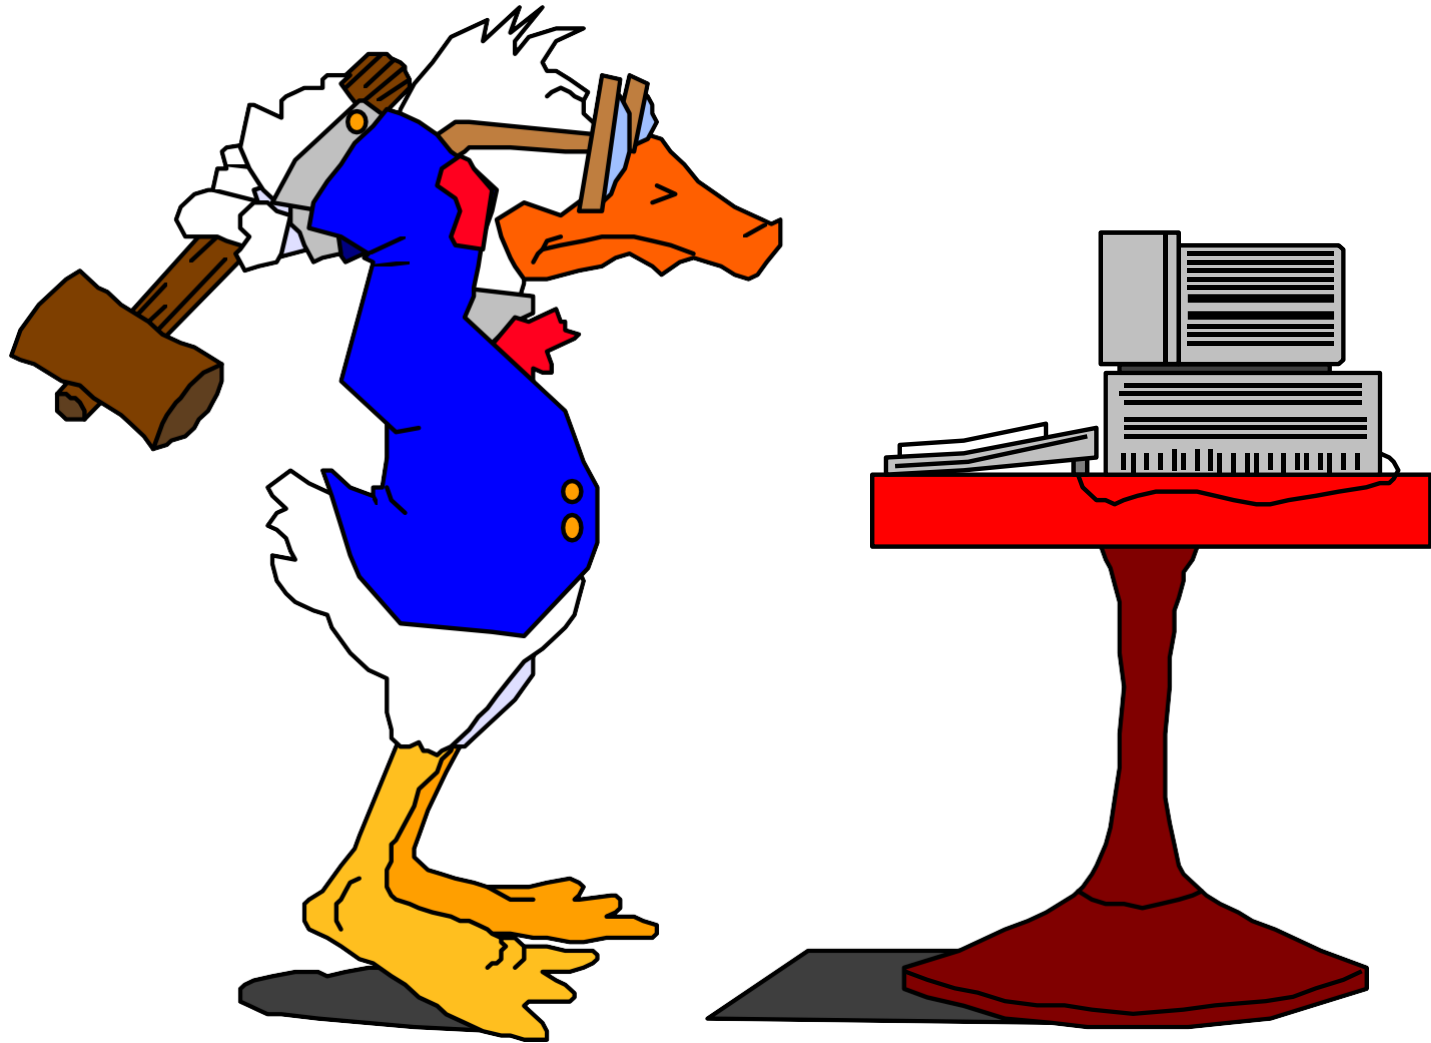
- Uniquely identified by
 - An internet address
 - An end-to-end protocol (e.G. TCP or UDP)
 - A port number
- **Two types of (TCP/IP) sockets**
- **Stream sockets (e.G. Uses TCP)**
 - Provide reliable byte-stream service
- **Datagram sockets (e.G. Uses UDP)**
 - Provide best-effort datagram service
 - Messages up to 65500 bytes
- Socket extend the conventional UNIX I/O facilities
 - File descriptors for network communication
 - Extended the read and write system calls



Sockets



Socket Programming



Client-Server communication

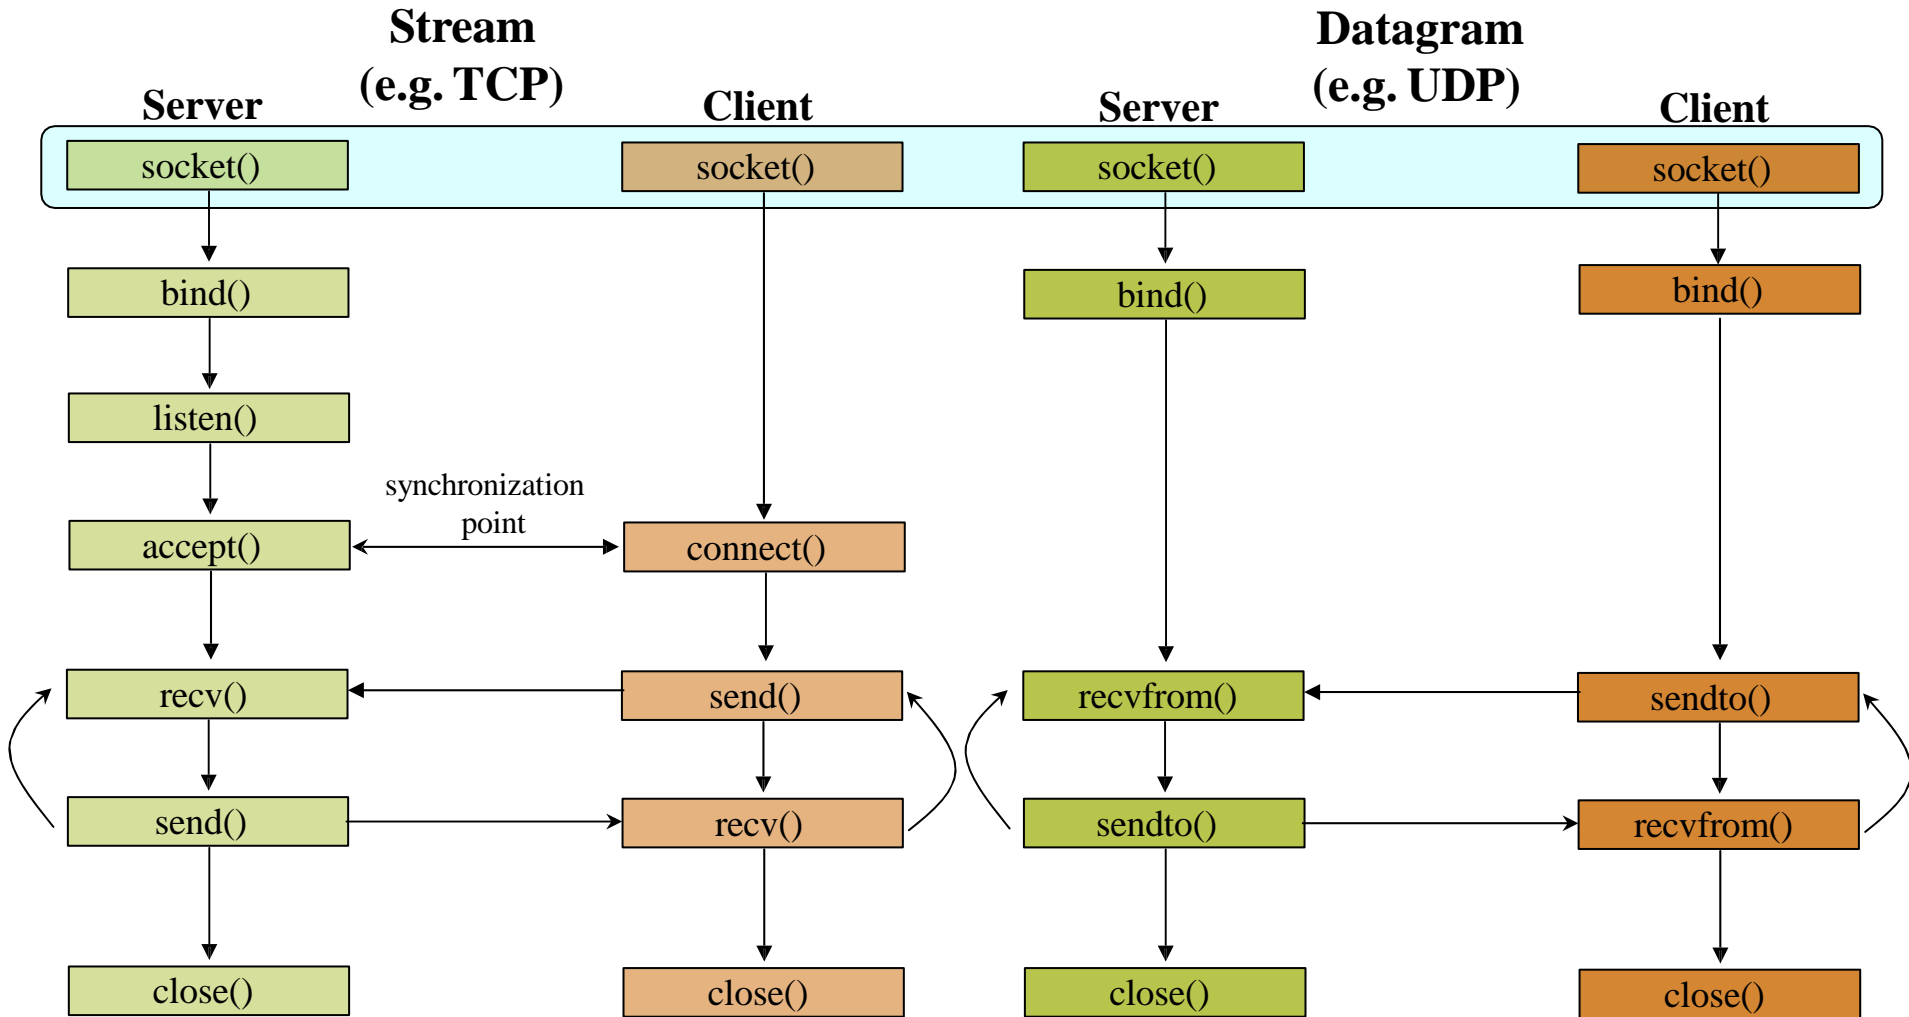
■ Server

- Passively waits for and responds to clients
- **Passive** socket

■ Client

- Initiates the communication
- Must know the address and the port of the server
- **Active** socket

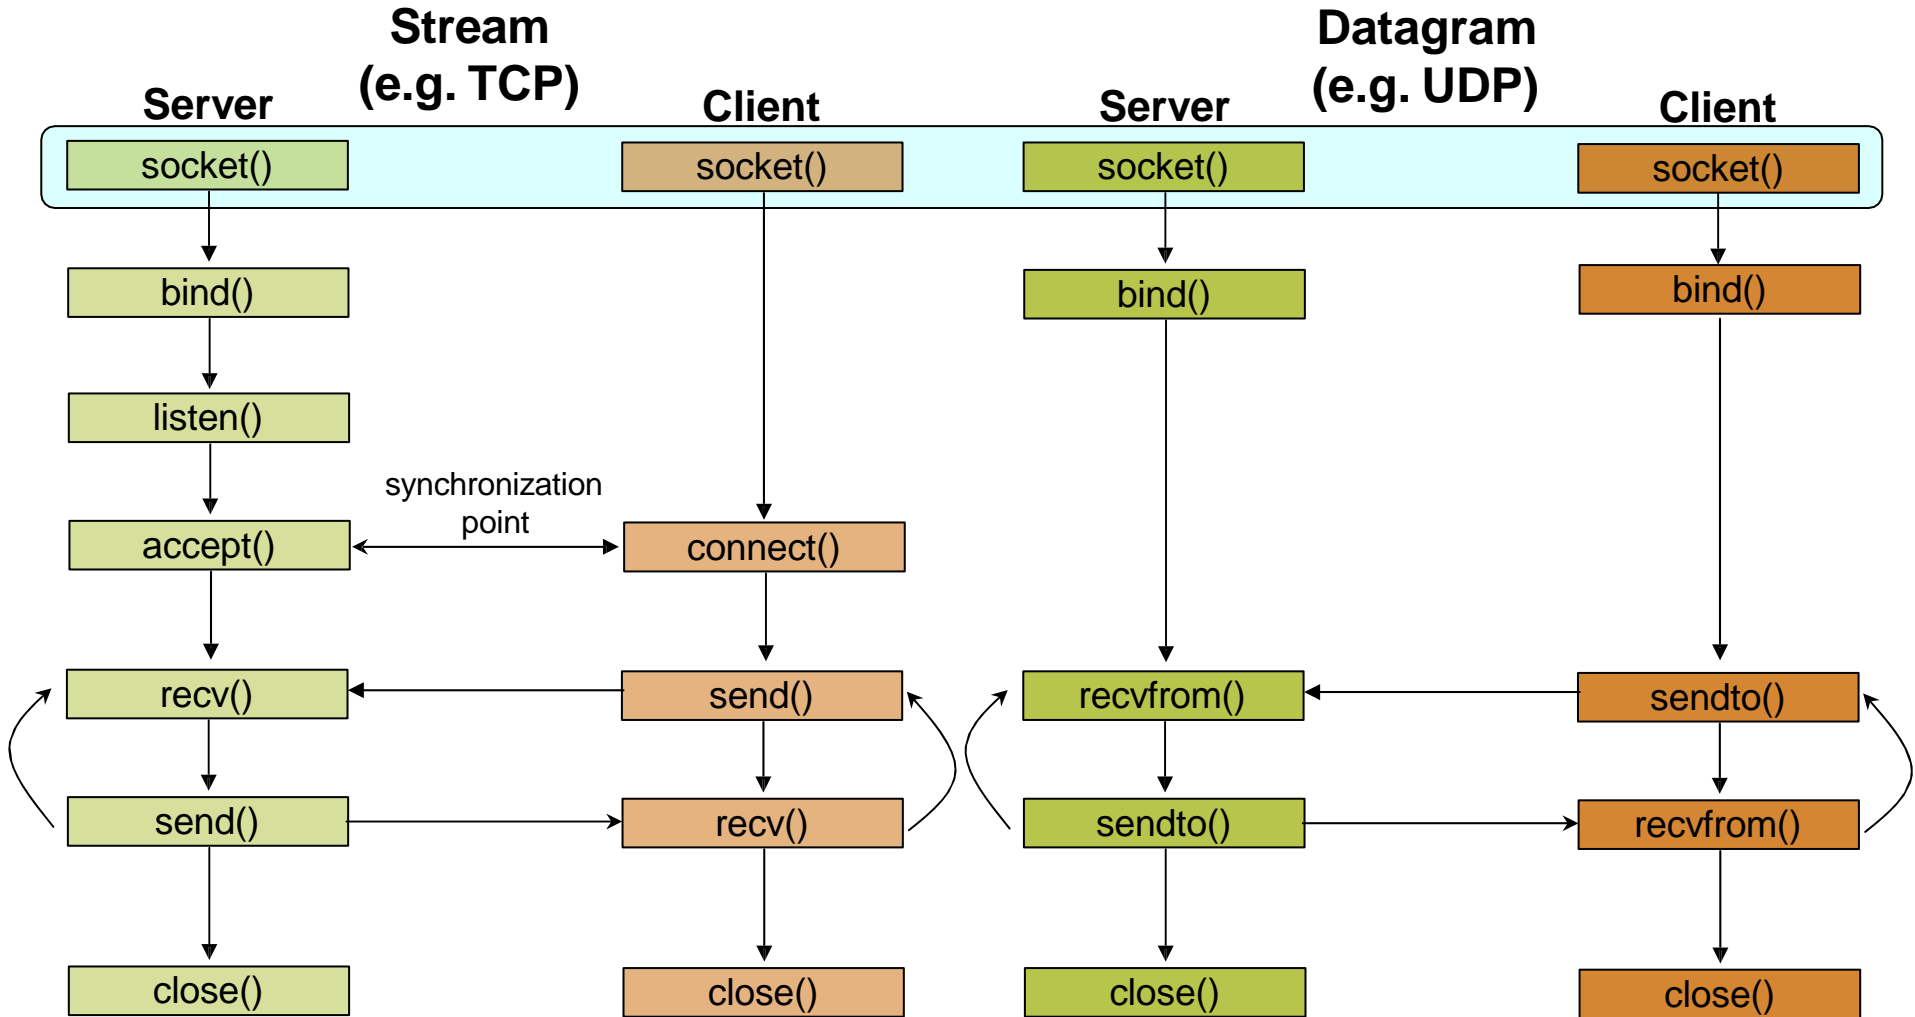
Client - Server Communication - Unix



Sockets - Procedures

Primitive	Meaning
Socket	Create a new communication endpoint
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections
Accept	Block caller until a connection request arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Receive	Receive some data over the connection
Close	Release the connection

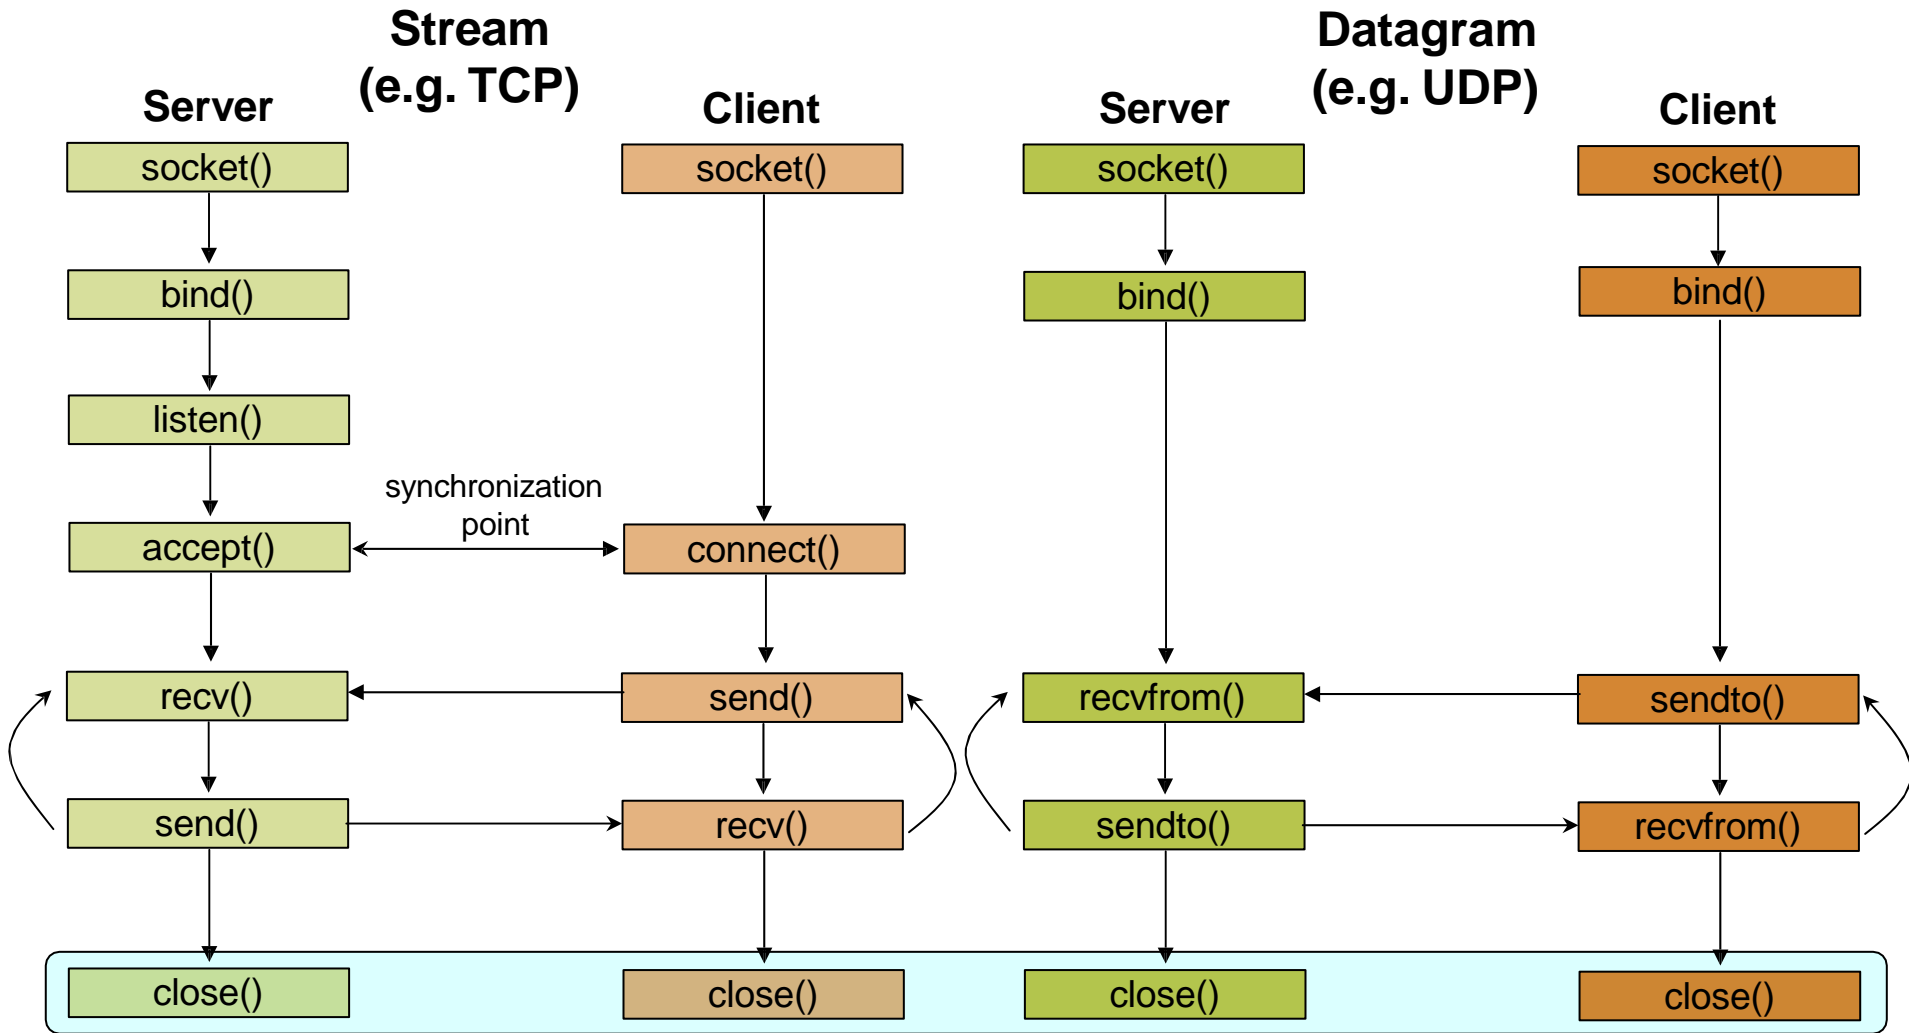
Client - Server Communication - Unix



Socket creation in C: `socket()`

- **`int sockid = socket(family, type, protocol);`**
 - **`sockid`**: socket descriptor, an integer (like a file-handle)
 - **`family`**: integer, communication domain, e.g.,
 - `PF_INET`, IPv4 protocols, Internet addresses (typically used)
 - `PF_UNIX`, Local communication, File addresses
 - **`type`**: communication type
 - `SOCK_STREAM` - reliable, 2-way, connection-based service
 - `SOCK_DGRAM` - unreliable, connectionless, messages of maximum length
 - **`protocol`**: specifies protocol
 - `IPPROTO_TCP` `IPPROTO_UDP`
 - usually set to 0 (i.e., use default protocol)
 - upon failure returns -1
- NOTE: socket call does not specify where data will be coming from, nor where it will be going to – it just creates the interface!

Client - Server Communication - Unix



Socket close in C: `close()`

- When finished using a socket, the socket should be closed
- `status = close(sockid);`
 - `sockid`: the file descriptor (socket being closed)
 - `status`: 0 if successful, -1 if error
- Closing a socket
 - closes a connection (for stream socket)
 - frees up the port used by the socket

Specifying Addresses

- Socket API defines a **generic** data type for addresses:

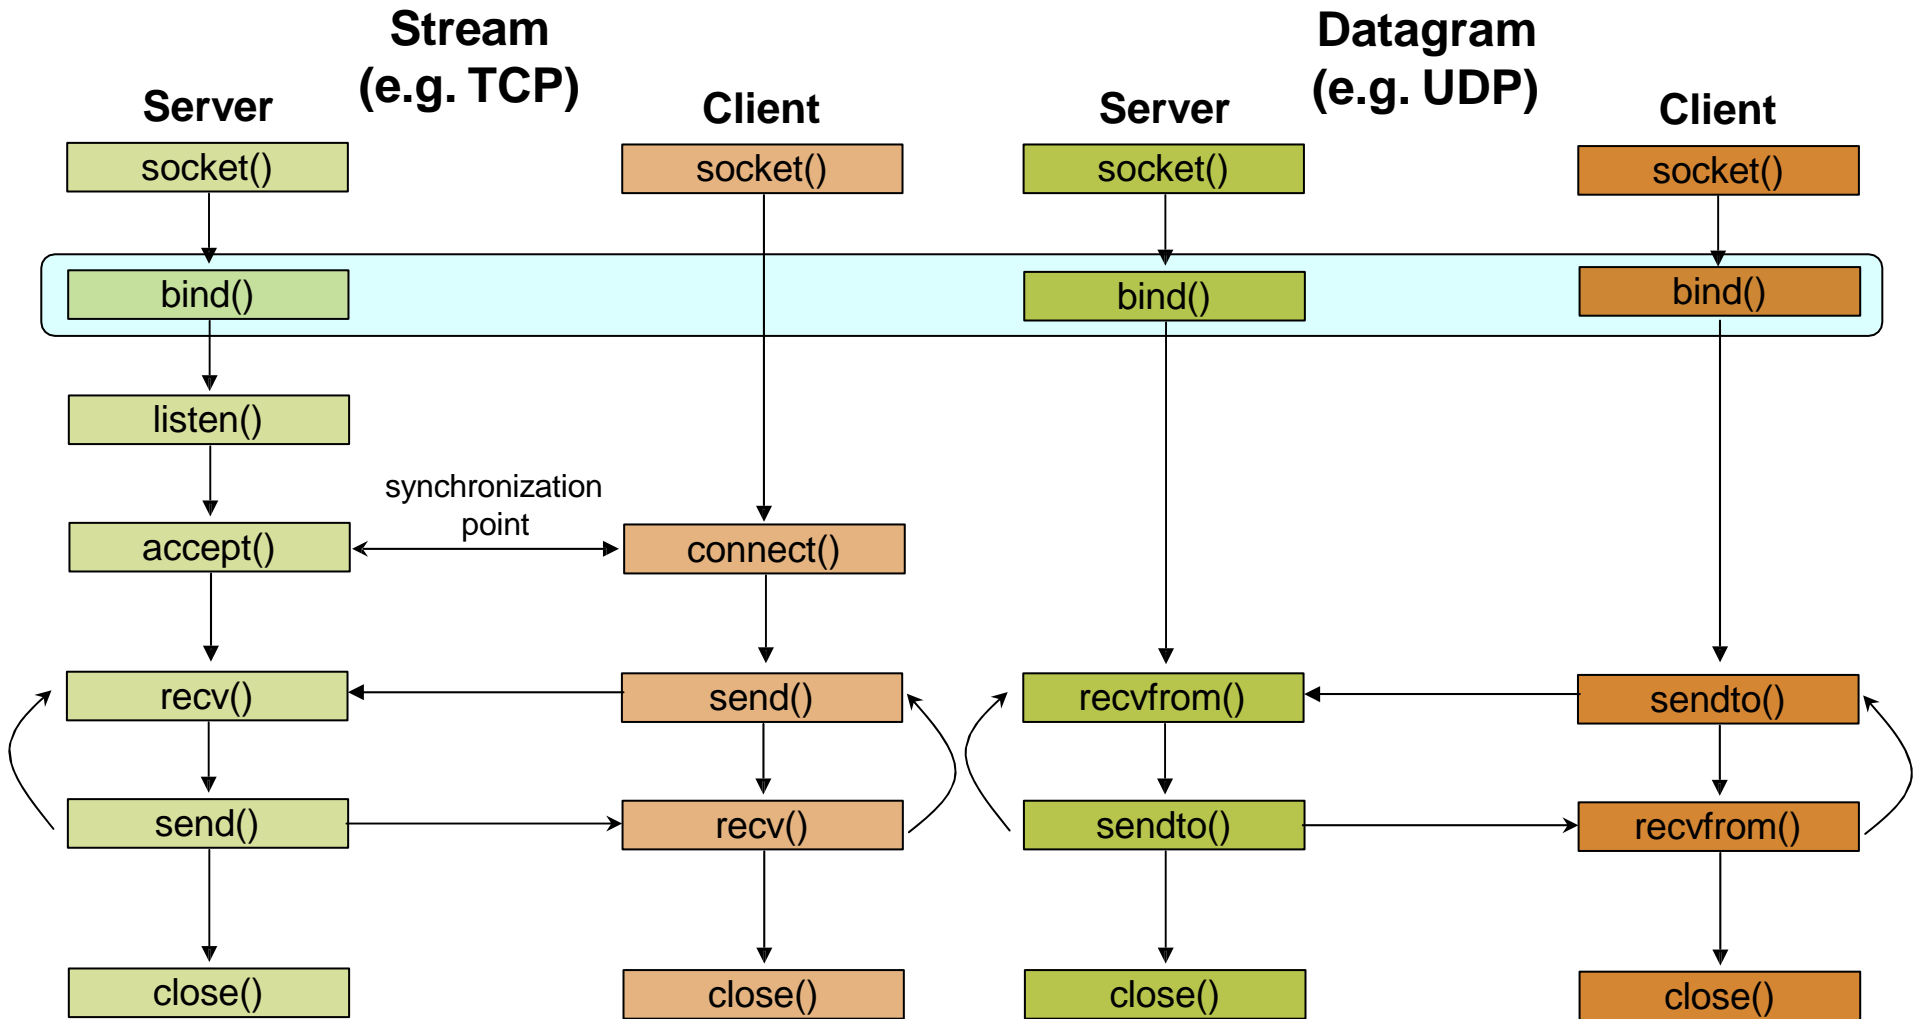
```
struct sockaddr {  
    unsigned short sa_family; /* Address family (e.g. AF_INET) */  
    char sa_data[14];        /* Family-specific address information */  
}
```

- Particular form of the sockaddr used for **TCP/IP** addresses:

```
struct in_addr {  
    unsigned long s_addr; /* Internet address (32 bits) */  
}  
  
struct sockaddr_in {  
    unsigned short sin_family; /* Internet protocol (AF_INET) */  
    unsigned short sin_port;   /* Address port (16 bits) */  
    struct in_addr sin_addr;   /* Internet address (32 bits) */  
    char sin_zero[8];         /* Not used */  
}
```

- **Important:** sockaddr_in can be casted to a sockaddr

Client - Server Communication - Unix



Assign address to socket: bind()

- associates and reserves a port for use by the socket
- **int status = bind(sockid, &addrport, size);**
 - **sockid**: integer, socket descriptor
 - **addrport**: struct sockaddr, the (IP) address and port of the machine
 - for TCP/IP server, internet address is usually set to INADDR_ANY, i.e., chooses any incoming interface
 - **size**: the size (in bytes) of the addrport structure
 - **status**: upon failure -1 is returned

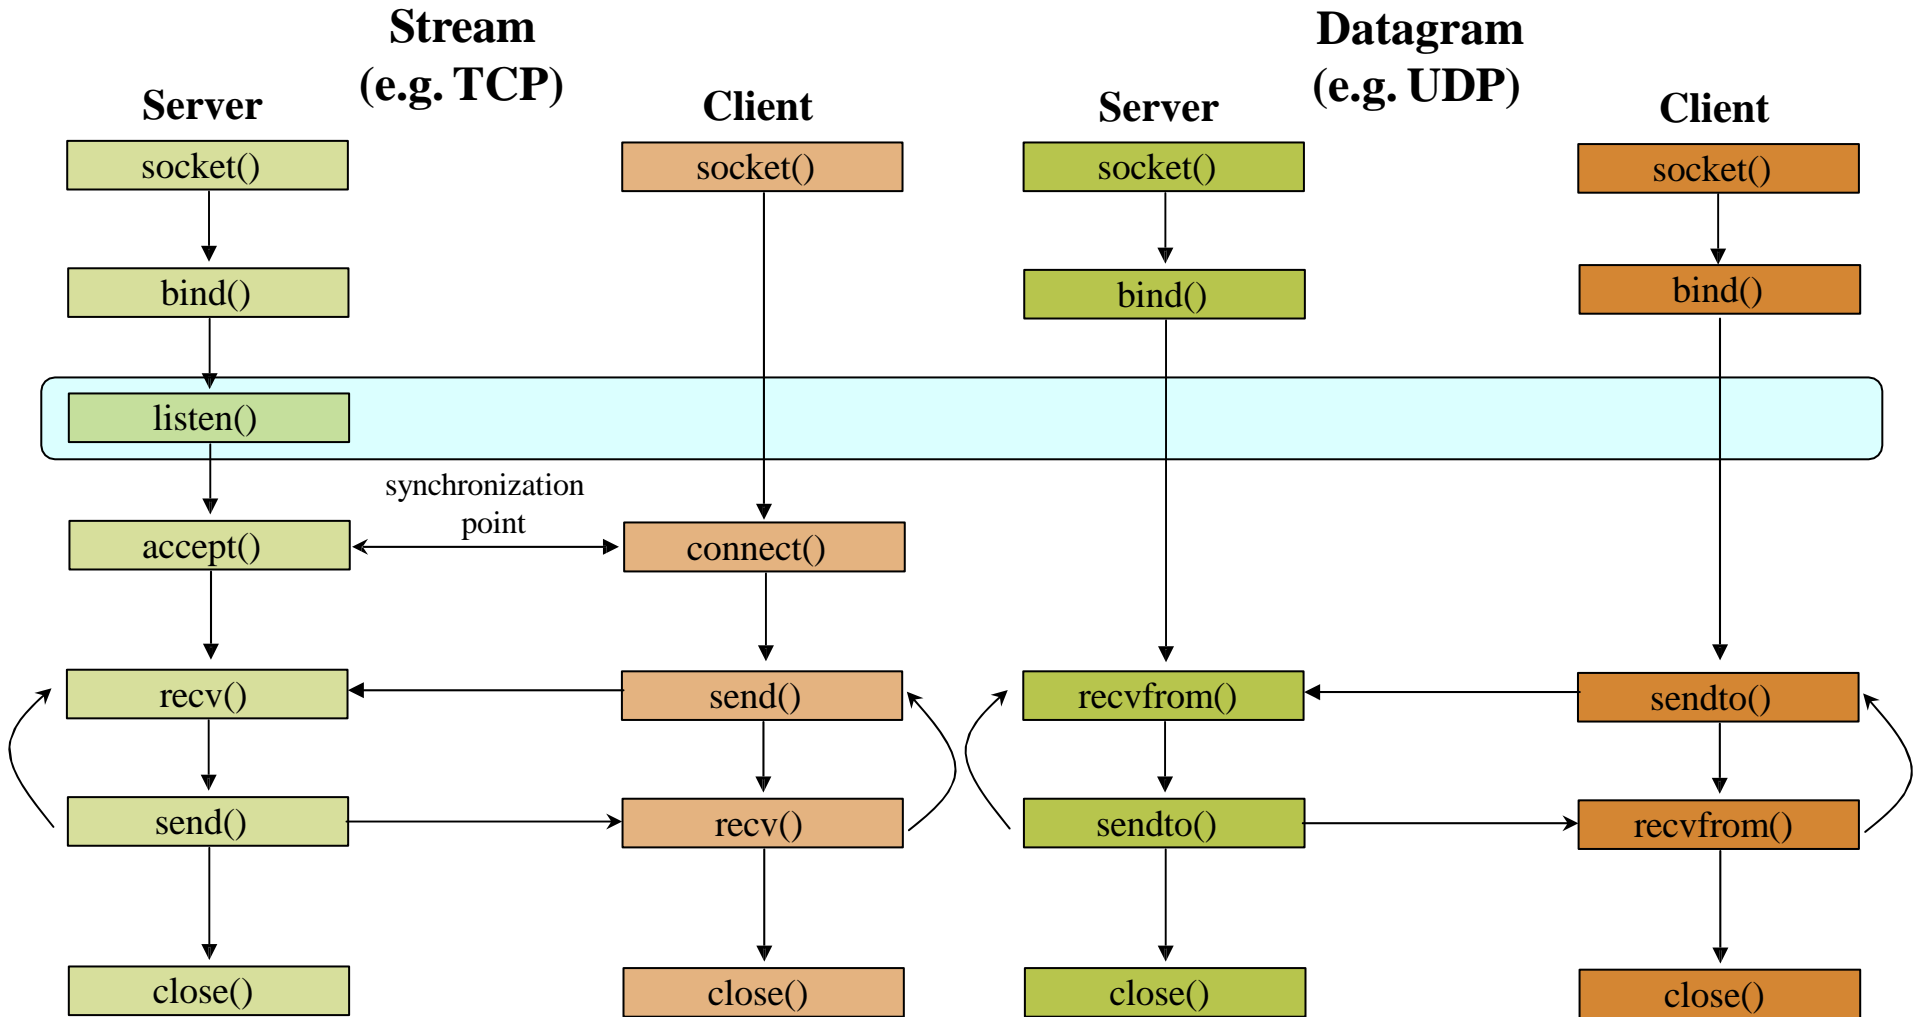
bind() - Example with TCP

```
int sockid;  
struct sockaddr_in addrport;  
sockid = socket(PF_INET, SOCK_STREAM, 0);  
  
addrport.sin_family = AF_INET;  
addrport.sin_port = htons(5100);  
addrport.sin_addr.s_addr = htonl(INADDR_ANY);  
if(bind(sockid, (struct sockaddr *) &addrport, sizeof(addrport)) != -1) {  
    ...  
}
```

Skipping the `bind()`

- Bind can be skipped for both types of sockets
- Datagram socket:
 - ❑ If only sending, no need to bind.
 - ❑ The OS finds a port each time the socket sends a packet
 - ❑ If receiving, need to bind
- Stream socket:
 - ❑ Destination determined during connection setup
 - ❑ Don't need to know port sending from (during connection setup, receiving end is informed of port)

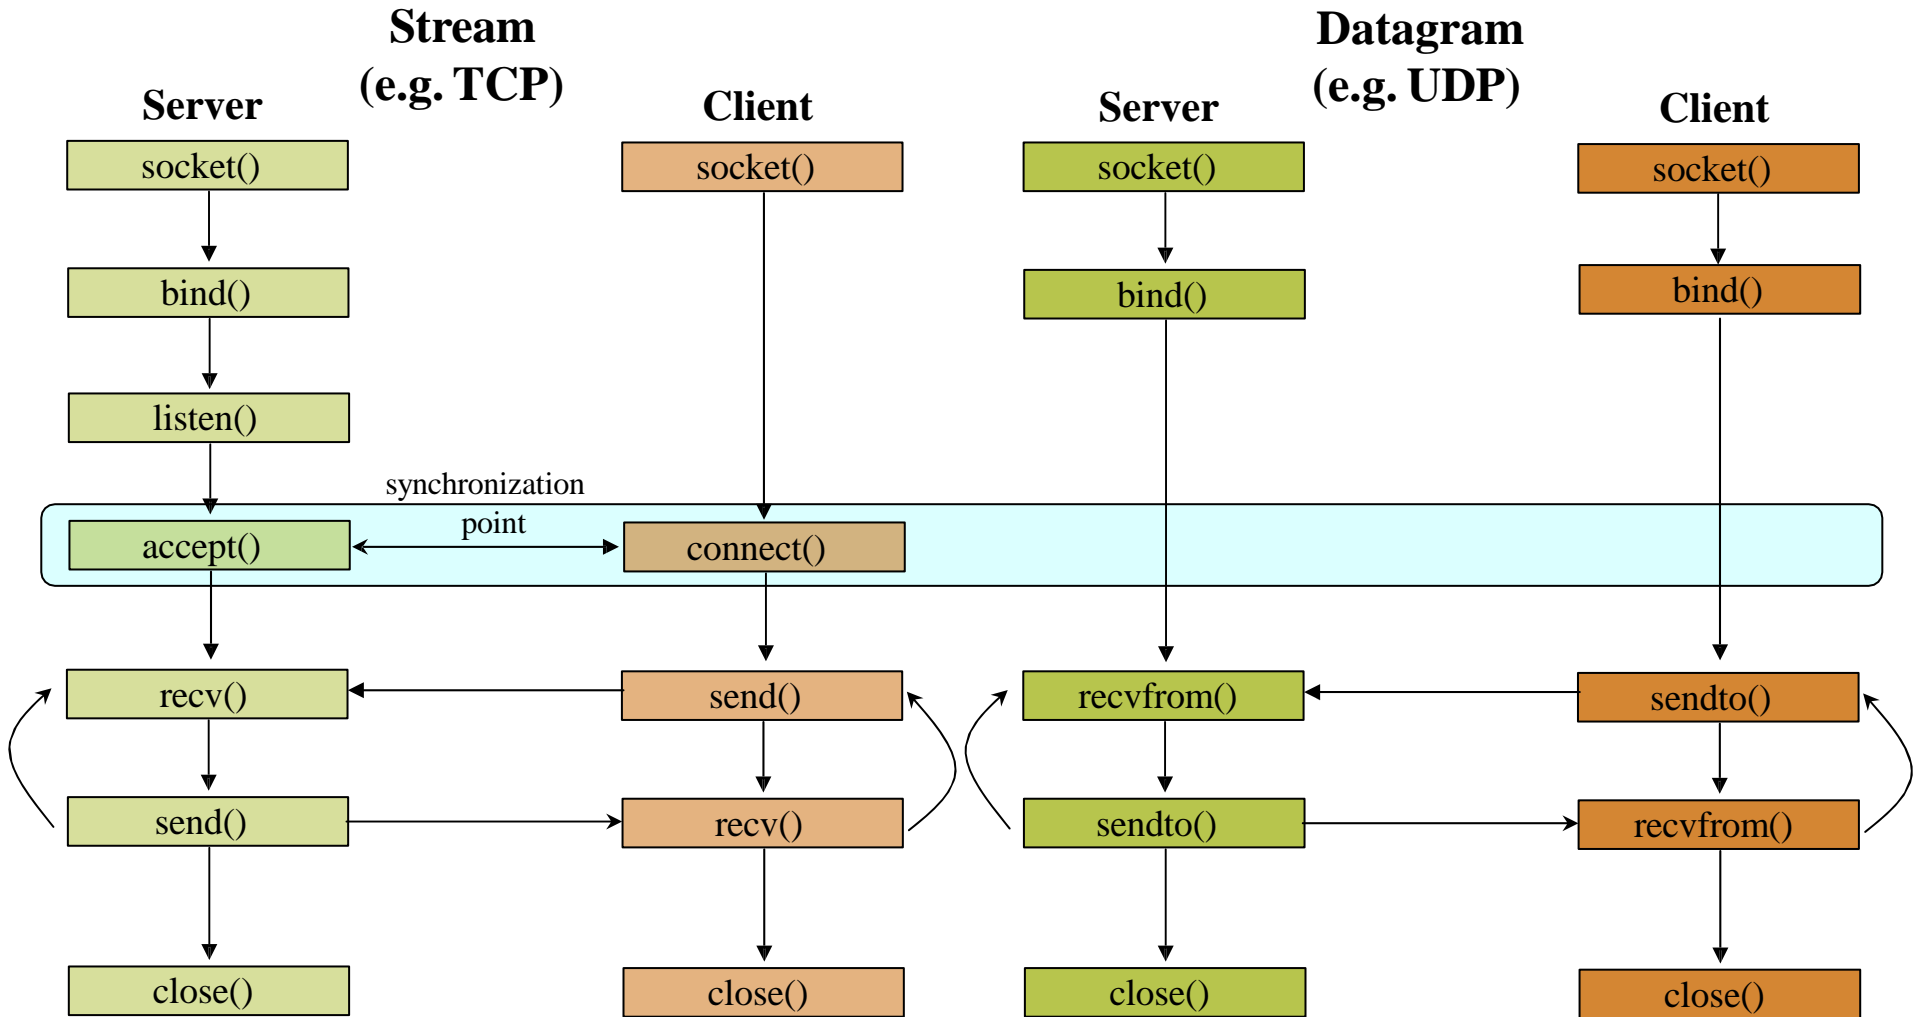
Client - Server Communication - Unix



Assign address to socket: bind()

- Instructs TCP protocol implementation to listen for connections
- **int status = listen(sockid, queueLimit);**
 - **sockid**: integer, socket descriptor
 - **queueLimit**: integer, # of active participants that can “wait” for a connection
 - **status**: 0 if listening, -1 if error
- listen() is **non-blocking**: returns immediately
- The listening socket (sockid)
 - is never used for sending and receiving
 - is used by the server only as a way to get new sockets

Client - Server Communication - Unix



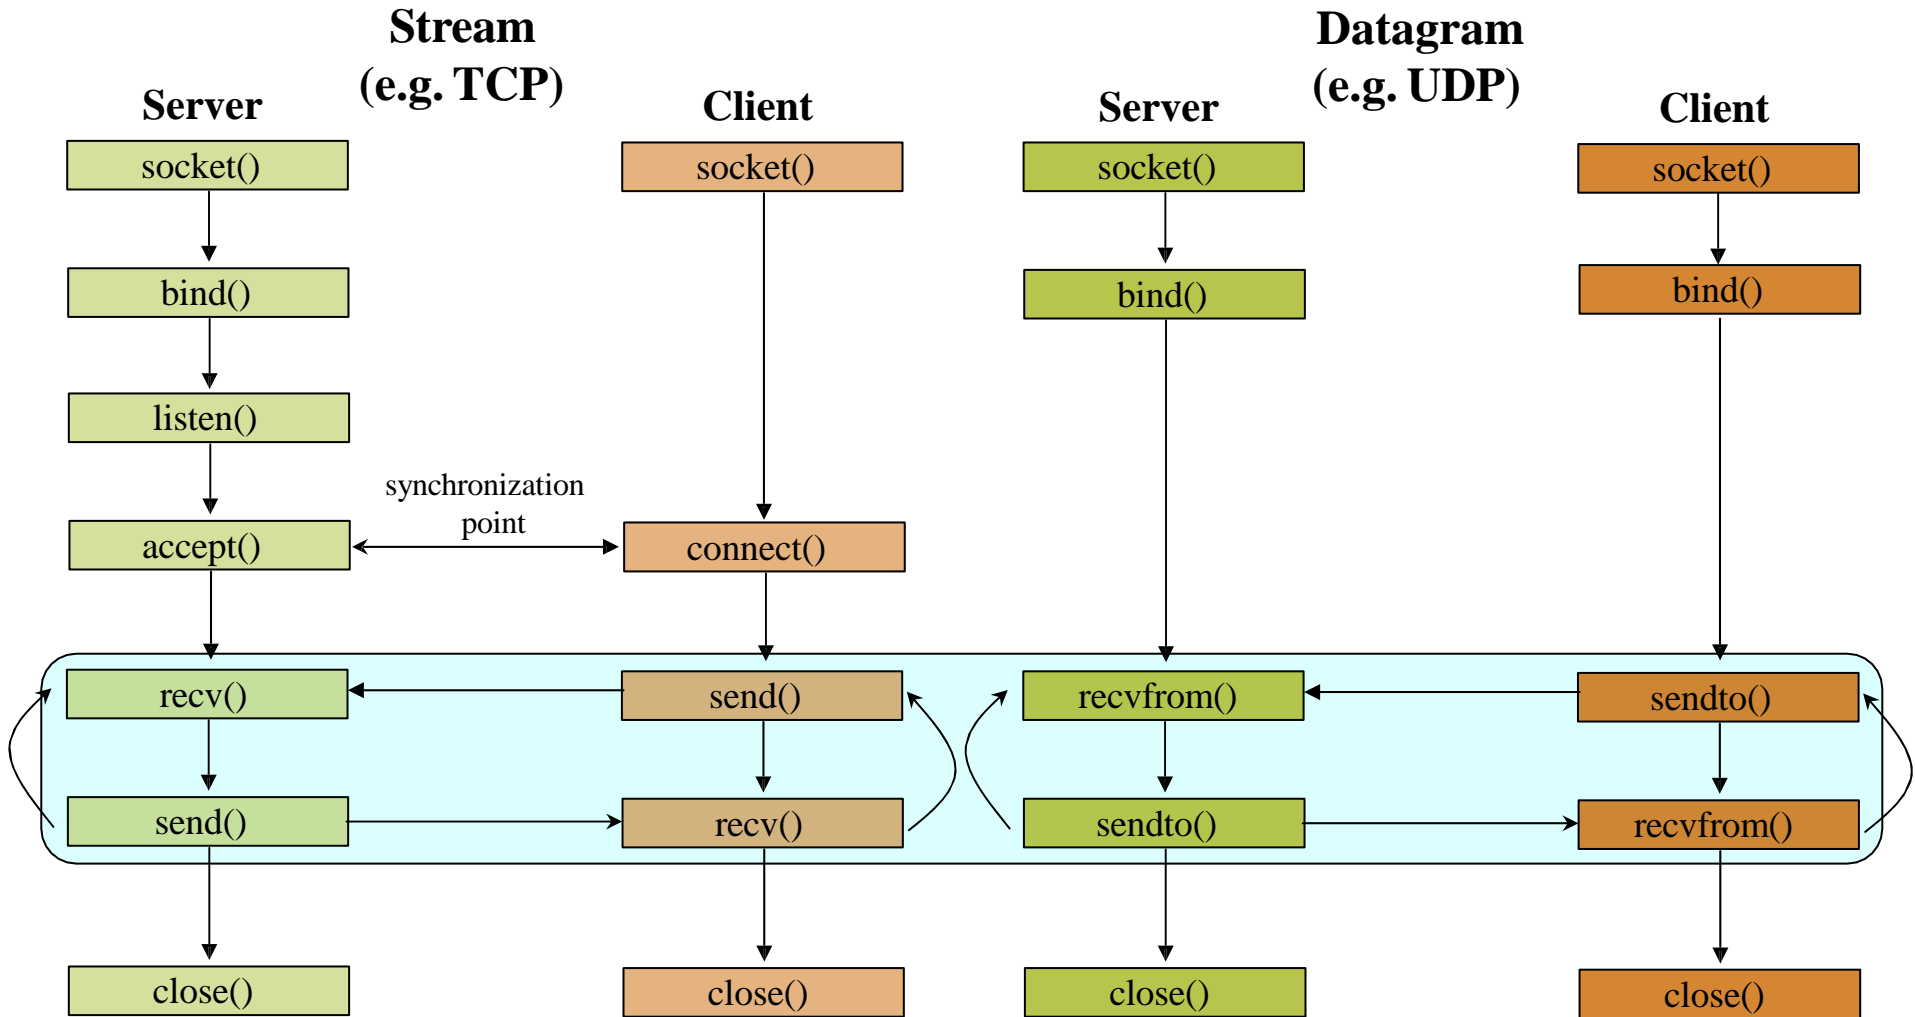
Establish Connection: connect()

- The client establishes a connection with the server by calling connect()
- **int status = connect(sockid, &foreignAddr, addrlen);**
 - **sockid**: integer, socket to be used in connection
 - **foreignAddr**: struct sockaddr: address of the passive participant
 - **addrlen**: integer, sizeof(name)
 - status: 0 if successful connect, -1 otherwise
- connect() is blocking

Incoming Connection: `accept()`

- The server gets a socket for an incoming client connection by calling `accept()`
- `int s = accept(sockid, &clientAddr, &addrLen);`
 - `s`: integer, the new socket (used for data-transfer)
 - `sockid`: integer, the orig. socket (being listened on)
 - `clientAddr`: struct `sockaddr`, address of the active participant
 - filled in upon return
 - `addrLen`: `sizeof(clientAddr)`: value/result parameter
 - must be set appropriately before call
 - adjusted upon return
- `accept()`
 - is **blocking**: waits for connection before returning
 - dequeues the next connection on the queue for socket (`sockid`)

Client - Server Communication - Unix



Exchanging data with stream socket

- **int count = send(sockid, msg, msgLen, flags);**
 - ❑ **msg**: const void[], message to be transmitted
 - ❑ **msgLen**: integer, length of message (in bytes) to transmit
 - ❑ **flags**: integer, special options, usually just 0
 - ❑ **count**: # bytes transmitted (-1 if error)
- **int count = recv(sockid, recvBuf, bufLen, flags);**
 - ❑ **recvBuf**: void[], stores received bytes
 - ❑ **bufLen**: # bytes received
 - ❑ **flags**: integer, special options, usually just 0
 - ❑ **count**: # bytes received (-1 if error)
- Calls are **blocking**
 - ❑ returns only after data is sent / received

Exchanging data with datagram socket

- `int count = sendto(sockid, msg, msgLen, flags, &foreignAddr, addrlen);`
 - `msg, msgLen, flags`, count: same with `send()`
 - `foreignAddr`: struct `sockaddr`, address of the destination
 - `addrlen`: `sizeof(foreignAddr)`
- `int count = recvfrom(sockid, recvBuf, bufLen, flags, &clientAddr, addrlen);`
 - `recvBuf, bufLen, flags`, count: same with `recv()`
 - `clientAddr`: struct `sockaddr`, address of the client
 - `addrlen`: `sizeof(clientAddr)`
- Calls are **blocking**
 - returns only after data is sent / received

Self Practice –Echo Server and Client

- A client communicates with an “echo” server
- The server simply echoes whatever it receives back to the client